

Just Build It!

Tips for Making MLOps and ML Engineering **Real**

Andy McMahon*

ML Engineering Lead

NatWest Group, Edinburgh, Scotland

MLOps Community Meetup

**speaking in a personal capacity*

Some Preamble


- My talk title is slightly tongue-in-cheek
- Objective: Let go of the fear of getting started
- Some phrases that can capture my sentiment:
 - Have a bias for action/execution is key
 - Bootstrap your capability
- In a nutshell - attempt what you are **not ready to do** and you will learn faster!



MLOps and MLEng



ML Eng: Getting your data science models into production by applying enough software engineering.



MLOps: The end-to-end lifecycle management of models and solutions with ML in them.

MLOps and ML Eng

Software Engineering



ML Engineering

Disciplines with tools, tech, processes and best practice - part of the *means*

DevOps





MLOps

Culture of blended responsibilities - the *ends*

MLOps as Induction

1. Get your first model into production ($n=0$). **This is how you learn ML Engineering.**
2. Get the $(n+1)$ th model into production given that the n th model is in production. **This is how you learn MLOps.**



Challenge: Analysis Paralysis


mlflow

 dataiku

 fiddler


neptune.ai


ALGORITHMIA



Kedro

 ALIBI


Kubeflow

truera

 comet


DataRobot


TensorBoard

 Arthur



So What Do We Do?

The Chasm

Idea



Production




MODEL
LIMBO!!!

The Chasm

Idea

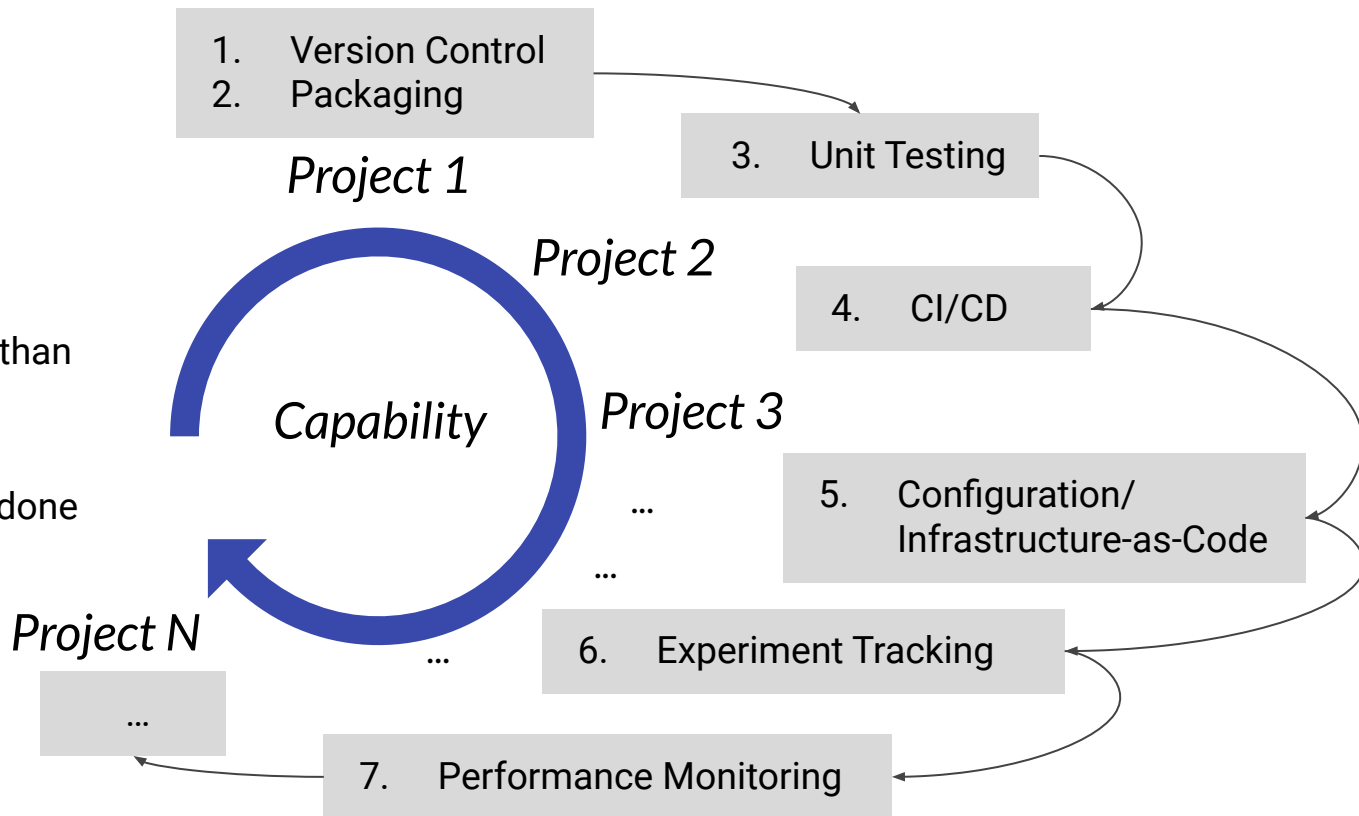


Production



Bootstrapping

- Start small then ratchet up
- Working is better than perfect
- You will never be done



Getting The (n=0)th Model Into Production

ML Eng 101 - Production

What?

- An isolated environment with strict controls, defined risk profile, service level agreements
- The place where your 'real' inference happens
- Where request/response traffic goes

How?

- Forget about notebooks* (sorry everyone, but it's true)
- We package, we lint, we test, we build, we version control
- We follow a clear route-to-live with appropriate promotion mechanisms

Done.

**We can fight about this one later ...*

ML Eng 102 - The Code

Prep for Prophet

```
df.rename(columns= {'Datetime': 'ds', 'AEP_MW': 'y'}, inplace=True)
```

```
df['ds']=df['ds'].astype('datetime64[ns]')
```

```
df.dtypes
```

```
#Initialize Split Class, we'll split our data 5 times for cv
ts_splits = TimeSeriesSplit(n_splits=5)
```

Train and Forecast

```
tmp = time_split_train_test(df.sort_values('ds', ascending=True).iloc[-1000:], ts_splits)
```

```
tmp.head()
```


Plot

```
nrow = 5; ncol = 1;
fig, axs = plt.subplots(nrows=nrow, ncols=ncol, figsize=(20,30))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
for i, ax in enumerate(fig.axes):
    split_rmse = tmp[(tmp['split']==i) & (tmp['train']==False)][['rmse']].iloc[0]
    ax.set_title('Split '+str(i)+' - RMSE: '+ '{:.2f}'.format(split_rmse))

    tmp[(tmp['split']==i) & (tmp['train']==True)].plot(x='ds', y='y', ax=ax, color='blue', marker='o')
    tmp[(tmp['split']==i) & (tmp['train']==False)].plot(x='ds', y='y', ax=ax, color='red', marker='o')
    tmp[(tmp['split']==i) & (tmp['train']==False)].plot(x='ds', y='yhat', ax=ax, color='orange', marker='^')
```

- OOP or Functional
- Separation of Concerns
- Keep it Simple Stupid
- Unit-test friendly


```
38 class Trainer(object):
39
40     def __init__(self, forecast_session):
41         self.forecast_session = forecast_session
42         self.df_pred_metadata = self.get_df_pred_metadata()
43         self.latest_predictor = self.get_latest_predictor()
44
45     def get_df_pred_metadata(self):
46         predictor_metadata = self.forecast_session.forecast.list_predictors()['PredictorArn']
47         df_pred_metadata = pd.DataFrame.from_records(predictor_metadata)
48         return df_pred_metadata
49
50     def get_latest_predictor(self):
51         latest_predictor = self.df_pred_metadata.sort_values(by='CreationTime', ascending=False).loc[0].to_dict()
52         return latest_predictor
53
54     def latest_predictor_in_tolerance(self, tolerance_days=2):
55         train_time_elapsed_days = (
56             datetime.datetime.now() - self.latest_predictor['CreationTime']).replace(tzinfo=None)
57             ).days
58         if train_time_elapsed_days < tolerance_days:
59             return True
60         else:
61             return False
62
63     def train_new_predictor(self):
64         PREDICTOR_NAME = PREDICTOR_BASE_NAME + datetime.datetime.now().strftime(format='%Y_%m_%d_%H_%M')
65         train_response = self.forecast_session.forecast.create_predictor(PredictorName=PREDICTOR_NAME,
66                                                                           AlgorithmArn=ALGORITHM_ARN,
67                                                                           ForecastHorizon=7,
68                                                                           PerformAutoML=False,
69                                                                           PerformHPQ=False,
70                                                                           InputDataConfig={
71                                                                               "DatasetGroupArn": DATASET_GROUP_ARN,
72                                                                           },
73                                                                           FeaturizationConfig={
74                                                                               "ForecastFrequency": DATASET_FREQUENCY
75                                                                           })
76         return train_response
77
78     def create_latest_forecast(self):
79         FORECAST_NAME = FORECAST_BASE_NAME + datetime.datetime.now().strftime(format='%Y_%m_%d_%H_%M')
80         create_forecast_response = self.forecast_session.forecast.create_forecast(
81             ForecastName=FORECAST_NAME,
82             PredictorArn=self.latest_predictor['PredictorArn'])
83         return create_forecast_response
```




ML Eng 103 - The Deployment

CI/CD Is Your Friend

- Github Actions, Jenkins, AWS CodePipeline/CodeBuild ...
- Automate, automate, automate!
- Checks and balances like minimum test coverage or even data quality check passes




```
1 # This is a basic workflow to help you get started with Actions
2
3 name: Upload DAGS to S3
4
5 # Controls when the action will run.
6 on:
7   # Triggers the workflow on push or pull request events but only for the main branch
8   push:
9     branches: [ main ]
10  pull_request:
11    branches: [ main ]
12
13 # Allows you to run this workflow manually from the Actions tab
14 workflow_dispatch:
15
16 jobs:
17   deploy:
18     name: Upload DAGS to Amazon S3
19     runs-on: ubuntu-latest
20
21     steps:
22     - name: Checkout
23       uses: actions/checkout@v2
24
25     - name: Configure AWS credentials from account
26       uses: aws-actions/configure-aws-credentials@v1
27       with:
28         aws-access-key-id: ${ secrets.AWS_ACCESS_KEY_ID }
29         aws-secret-access-key: ${ secrets.AWS_SECRET_ACCESS_KEY }
30         aws-region: us-east-1
31
32     - name: Copy files to bucket with the AWS CLI
33       run: |
34         aws s3 cp ./dags s3://github-actions-ci-cd-tests --recursive --include "*.py"
```



ML Eng 104 - The Deployment (Configuration)

Configuration Is Your Other Friend

- I love YAML!
- Separate out what is instance specific and what is generic application logic/code/modelling
- Reduce complexity of deployments needed for configuration changes!



Hydra
A framework for elegantly configuring complex applications

[Get Started](#) [Star](#) 5,311

```
@dataclass
class MySQLConfig:
    host: str = "localhost"
    port: int = 3306


@dataclass
class UserInterface:
    title: str = "My app"
    width: int = 1024
    height: int = 768

@dataclass
class MyConfig:
    db: MySQLConfig = MySQLConfig()
    ui: UserInterface = UserInterface()

cs = ConfigStore.instance()
cs.store(name="config", node=MyConfig)


@hydra.main(config_path=None, config_name="config")
def my_app(cfg: MyConfig) -> None:
    print(f"Title={cfg.ui.title}, size={cfg.ui.width}x{cfg.ui.height} pixels")

if __name__ == "__main__":
    my_app()
```



Getting The (n+1)th Model into Production

MLOps 101 - Model Management

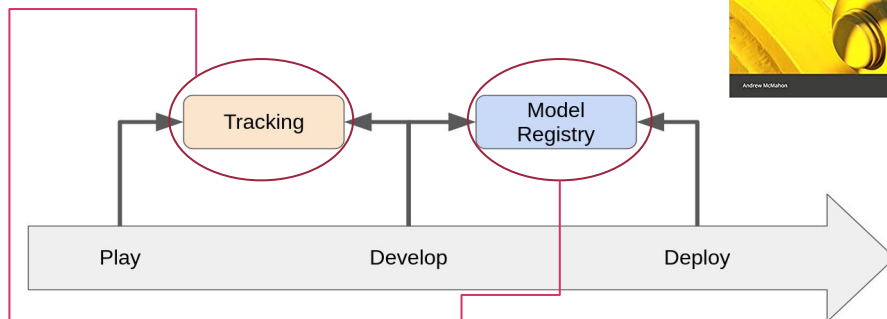


```
with mlflow.start_run(run_name="YOUR_RUN_NAME") as run:
    params = {
        'tol': 1e-2,
        'solver': 'sag'
    }
    # Fit a ridge classifier after performing standard scaling
    std_scale_clf = make_pipeline(StandardScaler(), RidgeClassifier(**params))
    std_scale_clf.fit(X_train, y_train)
    y_pred_std_scale = std_scale_clf.predict(X_test)

    mlflow.log_metrics(
        {
            'accuracy': metrics.accuracy_score(y_test, y_pred_std_scale),
            'precision': metrics.precision_score(y_test, y_pred_std_scale, average='macro'),
            'f1': metrics.f1_score(y_test, y_pred_std_scale, average='macro'),
            'recall': metrics.recall_score(y_test, y_pred_std_scale, average='macro')
        }
    )

    mlflow.log_params(params)



    # Log the sklearn model and register as version 1
    mlflow.sklearn.log_model(
        sk_model=std_scale_clf,
        artifact_path="sklearn-model",
        registered_model_name="sk-learn-std-scale-clf"
    )
```



```
# Transition the model stage to 'Staging'
client = MlflowClient()
client.transition_model_version_stage(
    name="sk-learn-std-scale-clf",
    version=1,
    stage="Staging"
)

# Transition the model stage to 'Production'
client = MlflowClient()
client.transition_model_version_stage(
    name="sk-learn-std-scale-clf",
    version=1,
    stage="Production"
)
```

MLOps 102 - Performance Monitoring



```
from sklearn.datasets import load_wine
from sklearn.model_selection import train_test_split



import alibi
from alibi_detect.cd import ChiSquareDrift, TabularDrift
from alibi_detect.utils.saving import save_detector, load_detector

# Grab the data
wine_data = load_wine()
feature_names = wine_data.feature_names
X, y = wine_data.data, wine_data.target

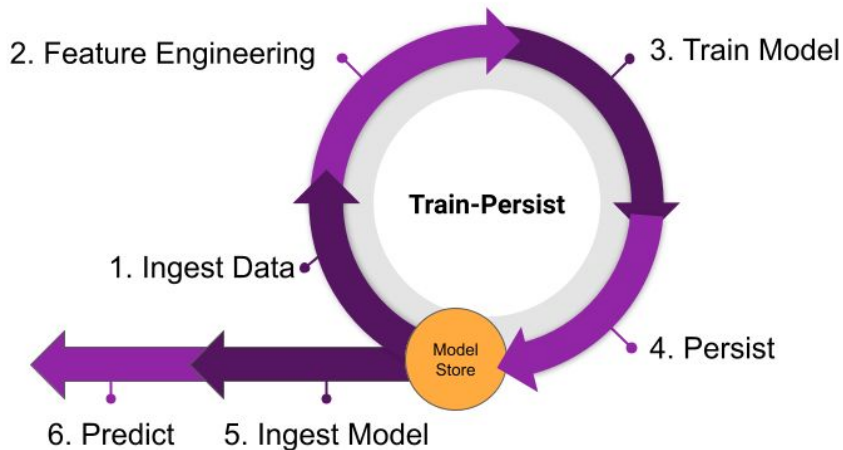
# Make a 50/50 reference/test split
X_ref, X_test, y_ref, y_test = train_test_split(X, y,
                                                test_size=0.50,
                                                random_state=42)

# Initialise the detector
cd = TabularDrift(p_val=.05, X_ref=X_ref)
# Check for drift
preds = cd.predict(X_test)
labels = ['No', 'Yes']
print('Drift: {}'.format(labels[preds['data']]['is_drift']))
```

MLOps 103 - Bringing it Together




The Train-Persist Process






Other Tips & Tricks: The 4P's

The 4 P's



People



Process



Product

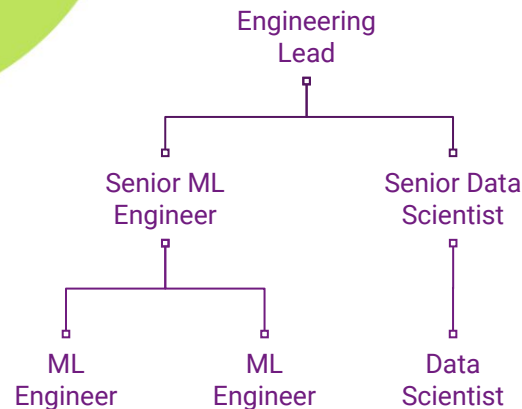
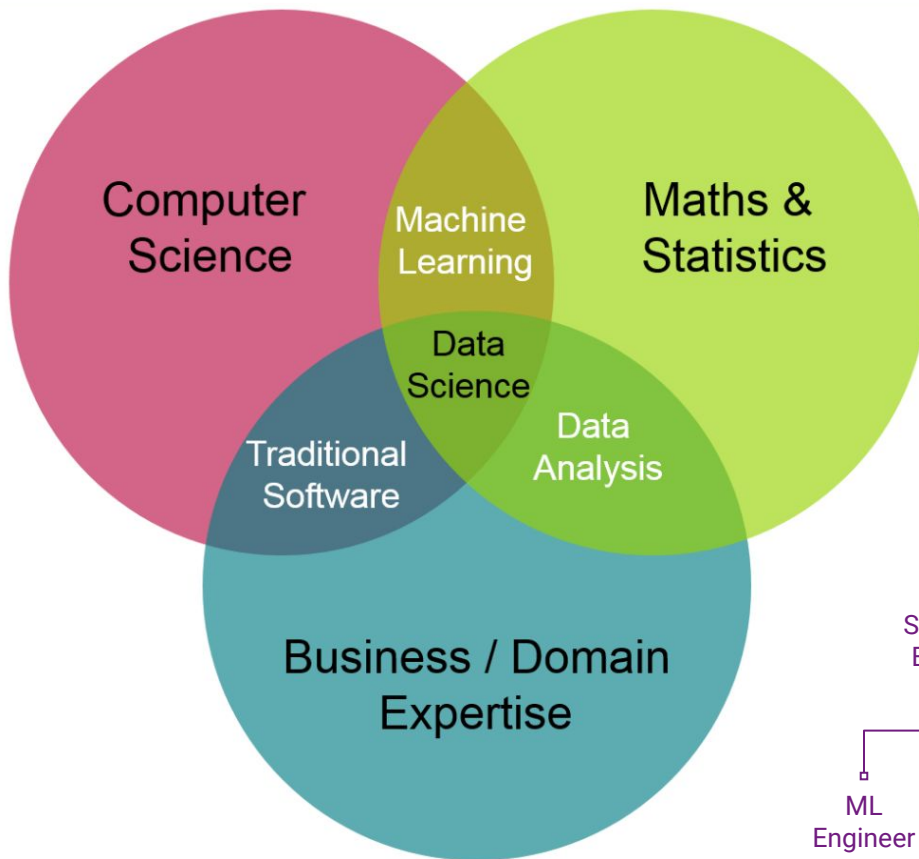



Pattern




People

People: Blended Teams > Unicorns



People: The Vision Should Be ...

Focussed



Relevant

Exciting



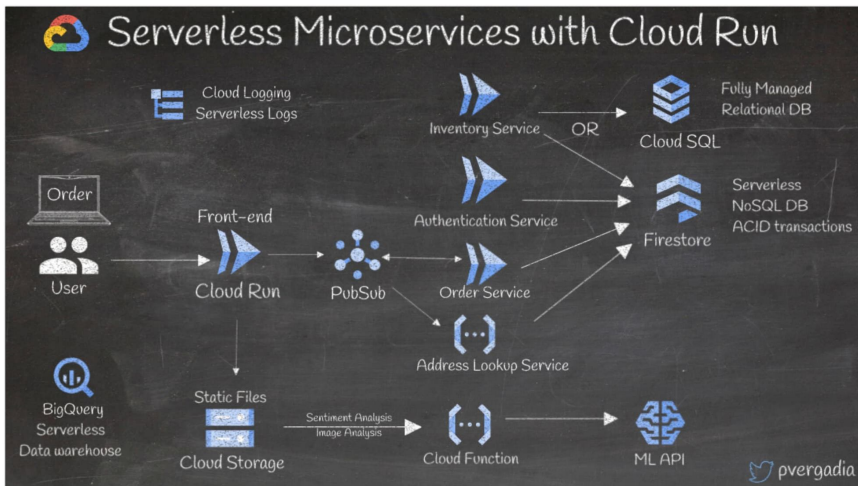
People: Know Your Customer






Patterns

Patterns: Reuse, Repeat, Recycle



13 sample architectures to kickstart your Google Cloud journey



Machine Learning Lens
AWS Well - Architected Framework


Patterns: Science != Engineering





Product

Products != Models != Analyses



MLEng/MLOps


[Hidden Technical Debt in Machine Learning Systems, Google](#)

Products work: Test, test, test ...

Products get delivered: Don't reinvent the wheel ...

Products come in a variety of shapes: Build for flexibility ...

Products != Models != Analyses



Process



Thank you!

Some places you can find me (say hello!)




@electricweegie



AndyMc629



Andy McMahon



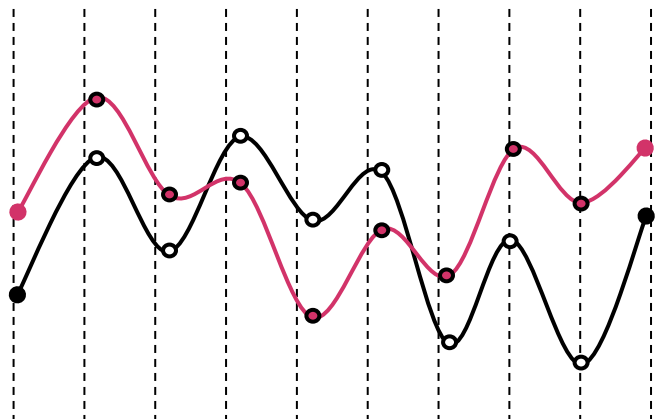
electricweegie.com





<https://shows.acast.com/ai-right>

Appendix




Decisions, decisions ...



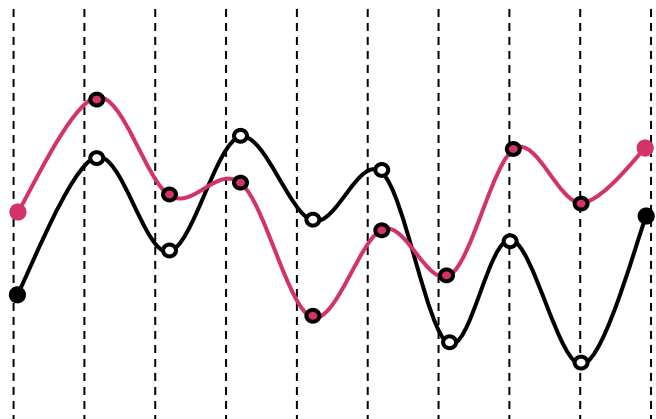
DATA





MODEL






Decisions, decisions ...



DATA



MODEL



Science != Engineering



ML Dev Process


Discovery

PoC

Development

Deployment

- Work with subject matter experts and stakeholders
- Understand challenge
- Define end goal (KPIs, baseline performance)
- Sketch MVP implementation
- Summarise for sign off



ML Dev Process


Discovery

PoC

Development

Deployment

- Exploratory data analysis
- Identification of 'no go' scenarios (blockers for the project)
- Identification of other project risks
- Bare bones implementation
- Evidence that the solution is viable and will solve the problem (sign off again)



databricks

ML Dev Process

Discovery

PoC

Development

Deployment

- Dev sprints
- Git strategy
- Unit and Integration tests
- Model version control
- Baseline performance tests/results



 PyTorch



databricks

 Spark
MLlib

 pytest

mlflow

 scikit
learn

 git



TensorFlow

ML Dev Process

Discovery

PoC

Development

Deployment

- Final packaging
- Deployment and update strategy
- Infrastructure instantiation
- Documentation completion
- Performance monitoring
- CI/CD pipeline definition

